


Candidate geometry × paired verification evidence

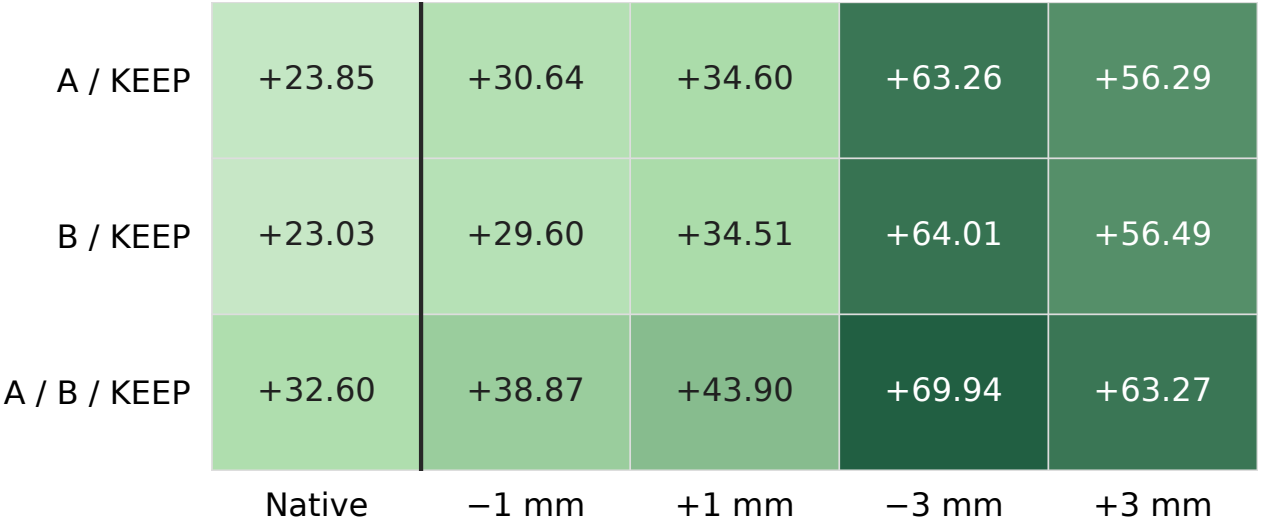
Predeclared primary: B_R balanced · 3 exposed replay scenes × 4 ROIs · ROI → scene equal weight

A Balanced selection: MSE reduction (%)



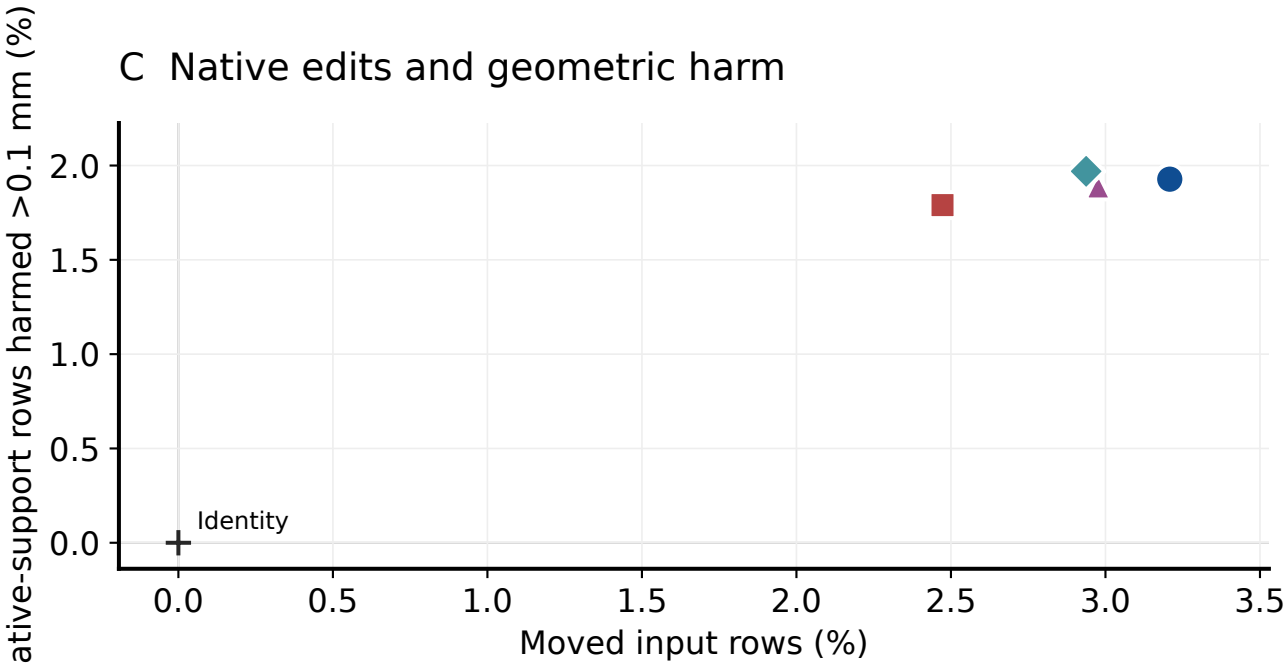
Native protection | Four repair conditions shown separately
Negative: higher error; positive: lower error versus identity

B Candidate potential: oracle reduction (%)



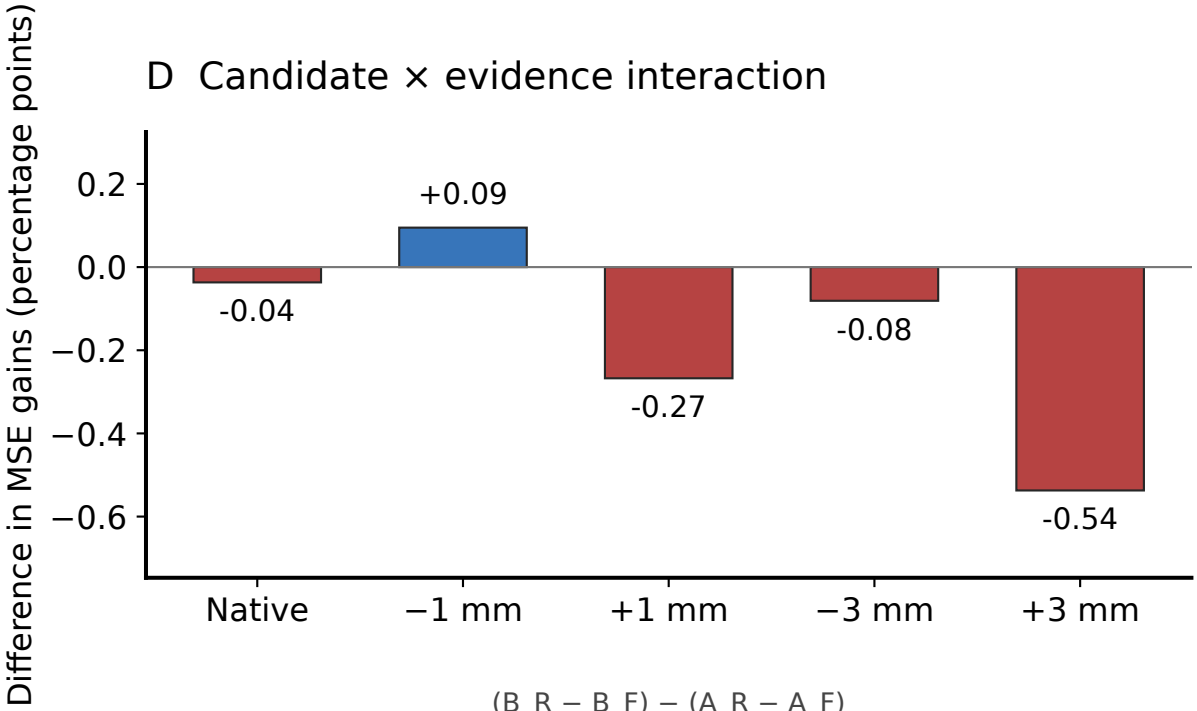
AB – A
gain (pp) +8.75 +8.23 +9.30 +6.69 +6.97

C Native edits and geometric harm



- A_F: moved 2.47%, harm 1.79%
- ▲ A_R: moved 2.98%, harm 1.89%
- ◆ B_F: moved 2.94%, harm 1.97%
- B_R: moved 3.21%, harm 1.93%

D Candidate × evidence interaction



$(B_R - B_F) - (A_R - A_F)$
Positive: reserved-source gain is larger for B than for A

A: single PCA / full patch; B: archived multi-normal / half patch. F: fitted sources; R: correction-reserved sources.

Panel B uses evaluator-only choices as fixed-candidate bounds; no oracle is deployed. Shared native support and identity denominator.

Descriptive exposed replay, not new confirmation. All five conditions retained; no point-count confidence intervals or per-condition winner selection.